

# Advanced NLP -

## Session 6: LLM Engineering

Prof. Dr. Richard Sieg  
TH Köln IWS - WS 25/26

Lecture	Date	Topic
1	03.12.2025	Introduction & NLP Recap
2	04.12.2025	RNNs and LSTMs
3	10.12.2025	Attentions & Transformers
4	11.12.2025	Transformer Based Models
5	17.12.2025	Hackathon / Check-In
6	18.12.2025	LLM Architecture
7	07.01.2026	LLM Engineering
8	08.01.2026	Hackathon / Check-In
9	14.01.2026	LLM Shortcomings
10	15.01.2026	Final Presentations


# Agenda

01. Tutorial: LLM Engineering


02. RAG

03. Evaluating LLM Applications


# The 3 Ingredients of LLMs



Process long sequences and context




Efficient training on huge datasets



Follow (human) instructions

Effectively

# Back to the Future




# Tutorial



How can we give an LLM  
access to data?


02.

Retrieval  
Augmented  
Generation  
(RAG)



# RAG: Main Idea

- Based on the user input, relevant documents are **retrieved** from a database
- Those documents are given to the LLM as additional context (**augmentation**)
- The LLM then **generates** the output based on the original input and retrieved documents



⚠ There are some frameworks that unify all of these steps but in the end I always ended up building it myself

## Step 1: Indexing

- Depending on the size of the documents and the context window of the LLM, we might need to chunk the documents first
  - Different chunking strategies: Number of characters, overlaps, headings etc
- Compute vector embeddings for each chunk using a “small” model
  - Popular choices: (Sentence)BERT, models by LLM providers
- Store those embeddings along with the original text and metadata in a vector database
  - Popular choices: qdrant, Weaviate, Chroma, Pinecone, Faiss, Cohere

## Step 2: Retrieval

- Use the LLM input as a query for the vector database
- Embed the query using the same model as for embedding the documents
- Using **Similarity Search**, retrieve the  $k$  most similar documents from the database, e.g. using **cosine similarity**

$$sim(query, document) = \frac{q \cdot d}{\|q\| \cdot \|d\|}$$

- This can be combined with standard retrieval techniques like keyword search

## Step 3: Generation

- The original input is augmented with the retrieved chunks and given to the LLM
- The retrieved chunks are indicated in the augmented prompt:  
`### CONTEXT or <context> ... </context>` (see tutorial)
- It is often a good idea to note in the system prompt that the LLM will receive context from a database

# Extended RAG Techniques

- **Indexing**


- Different chunking strategies, use different embeddings (full text, summary etc), metadata, hierarchical index

- **Retrieval**


- Rewrite query, query routing to different sources

- **Generation**

- Rerank retrieved sources, build summaries, context selection



# RAG Challenges



# RAG Challenges

- Key questions
  - *When and what should be retrieved? How is the context used?*
- Chunking strategy tradeoffs (size, overlap etc)
- “Internal Knowledge” of the LLM vs information in the database
- Robustness with incorrect information in the documents
- Quality and relevance of the context, context too noisy
- What is the best embedding model for the use case?
- Ever-growing context window of LLMs (“RAG is dead ”)
- “Lost in the middle” effect – information from the middle of a context is less likely used by the LLM

# Prompt Engineering or Fine-Tuning?

## Prompt Engineering

Tell the model what to do at runtime, every time  
“You are an expert in...”

### Benefits

- Adaptable – for many inputs
- Instant – no training
- Flexible – change anytime
- Experimenting is cheap

+ RAG

Give access to current knowledge

👉 You should start here

## Fine-Tuning

Modify model weights through training examples  
*Permanent behavior change*

### Only when

- Fixed & specific task
- A smaller model is needed:
- Lower latency
- Smaller costs

👉 Only in rare cases

- Inefficient
- Worse performance than fine-tuning
- Sensitivity to wording
- Lack of clarity

# Outlook: Other Important Techniques


- **Tool Use**
  - LLMs given access to a functions
- **Agents**
  - “LLM that runs tool in a loop to achieve a goal” – Simon Willison
- **Context Engineering**
  - Giving agents the correct context
- **Model Context Protocol (MCP)**
- **DSPy**
  - Automatic prompt optimization

03.

# Evaluating LLM Applications




# The Evaluation Challenge

- Two major evaluation paradigms
- **Discriminative**: Closed-ended evaluations
  - Classification, QA, extraction etc
  - Limited correct answers
  - Automatic evaluation possible
- **Generative**: Open-ended evaluations
  - Generation, conversations
  - Many valid answers
  - Evaluation is inherently difficult



High benchmark scores don't guarantee real-world performance. Models can exploit spurious correlations (e.g., negation words → contradiction) without truly understanding the task.

# Evaluating Text Generation



## Content Overlap Metrics (BLEU, ROUGE etc)

- Measure n-gram overlap with reference text
- Fast, cheap, and widely used
- Problem: no semantic understanding

Needs Reference Text

## Model Based Metrics (BERTScore, BLEURT etc)

- Use embeddings for semantic similarity
- Better correlation with human judgement


Needs Reference Text

## Human Evaluation

- Gold standard
- Expensive and slow

# LLM as a Judge

- Issue: Human annotation does not scale
- ***LLM as a Judge***
  - LLM receives input, output, and evaluation criteria
  - LLM produces judgement and explanation (critique)



## Common Pitfalls

- Position bias (prefers first answer)
- Length bias (prefers longer answer)
- Self-preference (LLM prefers own output)


## Best Practices

- Binary pass/fail decision
- Include critique
- Use task-specific criteria

# Production – Traces & Monitoring

- Logging & Traces
  - Input, output, latency, costs, prompts
  - Tools: Langfuse, Langsmith, many more


 **Rule Number 1**  
**Always look at your traces**



# Production Cycle

## Key Takeaways

- ✓ Look at your traces
- ✓ Start simple (binary pass/fail)
- ✓ Task-specific evals
- ✓ Validate validators
- ✓ Iterate continuously



You are the  
best judge of  
output quality.  
Don't rely on  
numbers